

Multiple Regression

José Pedro Conceição, Kiko Sánchez , Eloi Cirera

March 25, 2019

Table of Contents

Executive Summary

The first Objective was accurately Predicting Sales Volume, this has not been fully achieved because of the small data set provided to the team, we did train some models and made some predictions, however, they accurate and the errors are big, but it was the best information value we could extract from this sample. Predicting the sales of new Products using a reduced sample won't be accurate, besides being statistically unsound. We found out that what actually predicts success in volume are both, however the best predictor for success comes from service reviews with an importance of 100% according to a random forest algorithm, followed by a 50% importance of 4 star reviews.

Why 4 star reviews and not other reviews ? Well because they had to be taken out of our model training, they had levels of relationship with our predictor (Volume) so high that they were biasing the whole model, making the predictions even more unreliable. For example the 5 star review had a perfect correlation with the Volume, this means that the volume would grow at the same rate as the 5 star reviews increased, which does not translate into reality. Nonetheless here are our final predictions.

Number

Game Console

Game Console

Tablet

Tablet

NoteBook

Technical Report

Pre-process

We always start by assessing the importance of each variable, we achieve this by doing a correlation matrix and training a simple model, followed by a `varImp()`, which will give us in percentage the importance of the variable for the model's prediction. (We need to first use a correlation matrix to see the correlation values, and take out anything that might bias our model, otherwise the "biased", features will just appear at the top of the `varImp()` output).

We created a function to dummyfy the variables and to check if there was any NA values (in any attribute), and if they exist, remove them, I also included a function that removes outliers, and a function to subset the data into different product types.

Process functions

Pre-process

```
PPfunction <- function(data) {  
  N <- dummyVars(" ~ .", data = data)  
  N <- data.frame(predict(N, newdata = data))  
  N <- N[,colSums(is.na(N)) == 0]  
  N  
}
```

Remove Outliers

```
RmOut <- function(D,V)  
{  
  Out <- boxplot(D$V ,plot = FALSE)$out  
  K <- D[-which(D$V %in% Out),]
```

```
K  
}
```

Sub-set by product types

```
SubSetDataProductTypes <- function(data,p,p1 = 0,p2 = 0 , p3 =  
0 , p4 = 0)  
  
  {  
    if ( p1 == 0 && p2 == 0 && p3 == 0 && p4 == 0)  
    {  
      Nsub <- subset(data, data$ProductType == p)  
  
      return(Nsub)  
    }  
    else if (p2 == 0 && p3 == 0 && p4 == 0){  
      Nsub <- subset(data, data$ProductType == p)  
  
      Nsub2 <-subset(data, data$ProductType == p1)  
  
      Nsub2 <- rbind(Nsub,Nsub2)  
  
      return(Nsub2)  
    }  
    else if (p3 == 0 && p4 == 0)  
    {  
  
      Nsub <- subset(data, data$ProductType == p)  
  
      Nsub2 <-subset(data, data$ProductType == p1)  
  
      Nsub3 <- subset(data,data$ProductType == p2)  
  
      Nsub3 <- rbind(Nsub,Nsub2,Nsub3)  
  
      return(Nsub3)  
    }  
  
    else if (p4 == 0){  
      Nsub <- subset(data, data$ProductType == p)  
  
      Nsub2 <-subset(data, data$ProductType == p1)  
  
      Nsub3 <- subset(data,data$ProductType == p2)
```

```

    Nsub4 <- subset(data,data$ProductType == p3)
    Nsub4 <- rbind(Nsub,Nsub2,Nsub3,Nsub4)
    return(Nsub4)
  }
  else{
    Nsub <- subset(data, data$ProductType == p)
    Nsub2 <-subset(data, data$ProductType == p1)
    Nsub3 <- subset(data,data$ProductType == p2)
    Nsub4 <- subset(data,data$ProductType == p3)
    Nsub5 <- subset(data,data$ProductType == p4)
    Nsub5 <- rbind(Nsub,Nsub2,Nsub3,Nsub4,Nsub5)
    return(Nsub5)
  }
}

```

I know it's not the most pretty or effective way to do this, but it works.

Correlation Matrix :


```

EP <- PPfunction(EP)
EP <- RmOut(EP,Volume)

corr_all<-cor(EP)

corrplot::
corrplot(corr_all,type="upper",tl.pos="td",method="circle",tl.cex
= 0.5,tl.col='black',diag=FALSE)

```



This information does not differ from the module 1's counterpart, it's obvious because we are using the same data set.

We trained a random forest followed by the use of `varImp()` function that assess the importance of each variables (without the ones we took out by looking at the correlation matrix). But first we need to create Test and Training sets, we also came up with a simple function to automate the process.

Train and Test Set function

```
TrainAndTestSets <- function(label,p,data,seed){
  set.seed(seed)

  inTrain <- createDataPartition(y= label, p = p , list = FALSE)
  training <- data[inTrain,]
  testing <- data[-inTrain,]

  list(trainingSet=training,testingSet = testing)
}
```

```
EP <-
EP[,c(1,2,3,4,5,6,7,8,9,10,11,12,13,14,16,18,20,21,22,23,24,25,26
,27,28)]
```

```
List <- TrainAndTestSets(EP$Volume,0.75,EP,123)
```

```
fitcontrol <- trainControl(method = "repeatedcv", repeats = 4)

Model <- train(Volume~., data = EP,method = "rf", trcontrol =
fitcontrol , tunelenght = 5
, preProcess = c("center",
"scale"),importance=T)

varImp(Model)
```

```
## rf variable importance
##
## only 20 most important variables shown (out of 24)
##
```

	Overall
## PositiveServiceReview	100.000
## x4StarReviews	46.607
## ProductWidth	11.786
## ShippingWeight	10.171
## Price	9.865
## x2StarReviews	9.387
## ProductType.ExtendedWarranty	8.821
## ProductType.Printer	7.122
## ProfitMargin	6.530
## ProductDepth	6.083
## ProductType.Tablet	6.030
## ProductType.Software	5.748
## ProductHeight	4.535
## ProductType.GameConsole	3.915
## Recommendproduct	3.785
## ProductType.Accessories	3.465
## ProductType.Smartphone	3.014
## ProductType.PC	2.530
## ProductType.Display	2.499
## NegativeServiceReview	2.344

So the only variables with a significant impact are only PostiveServiceReview and x4StarReviews.

Models and Predictions

Training Function

I created a function that trains every different model, by user specification

```
EP <- read.csv( file
="/home/zordo/Documents/Ubiquim/R-M2Task3/data/Epa.csv" , header =
TRUE , sep = ',')

EP <- EP[,c(1,5,9,18)]

EP <- PPfunction(EP)

EP <- RmOut(EP)

List <- TrainAndTestSets(EP$Volume,0.75,EP,123)

#### Random Forest ####
ModelRandomForest <-
TrainingFunction("rf",Volume~.,List$trainingSet,5)

PredictionRandomForest <-
predict(ModelRandomForest,List$testingSet)

TestResultsRF <-
postResample(PredictionRandomForest,List$testingSet$Volume)

#### SVM ####

svm.model <-
TrainingFunction("svm",Volume~.,List$trainingSet,5,10000000,0.000
0001)

svm.pred <- predict(svm.model,List$testingSet)

TestResultsSVM <- postResample(svm.pred,List$testingSet$Volume)

#### knn ####

KNN <- TrainingFunction("knn",Volume~.,List$trainingSet,30)
```

```

KnnPrediction <- predict(KNN,List$testingSet)

TestResultsKNN <-
postResample(KnnPrediction,List$testingSet$Volume)

####

AllTestResults <-
cbind(TestResultsKNN,TestResultsRF,TestResultsSVM)

AllTestResults

##           TestResultsKNN TestResultsRF TestResultsSVM
## RMSE           429.1288137    271.9646084    788.7251446
## Rsquared        0.6828584      0.8447768      0.2698622
## MAE            250.9259259    156.6409058    553.4143148

```

And then did another one to train the three models at the same time with a for loop

```

TrainAll3Models <- function (formula,data)
{

  Model <- vector(mode="list", length=length(methods))

  methods <- c("rf","svm","knn")

  for(i in 1:length(methods))
  {

    Model[[i]] <-
TrainingFunction(methods[i],formula,data,5)

  }
  Model
}

```

I didn't use this function that much since the mentors showed us another way of training without any function, and it's much easier and cleaner.

```

a <- c("Volume ~ x4StarReviews","Volume ~.", "Volume ~
PositiveServiceReview")

```



```

b <- c("lm", "rf", "knn", "svmLinear")
compare_var_mod <- c()

for (i in a) {
  for (j in b) {

    model <- train(formula(i), data = List$trainingSet, method =
b, trainControl=trainControl(method = "repeatedcv", repeats = 4))

    pred <- predict(model, newdata = List$testingSet)

    pred_metric <- postResample(List$testingSet$Volume, pred)

    compare_var_mod <- cbind(compare_var_mod , pred_metric)

  }
}

compare_var_mod

##          pred_metric pred_metric pred_metric pred_metric
pred_metric
## RMSE      459.4407974 459.4407974 459.4407974 459.4407974
482.8296811
## Rsquared    0.5286655    0.5286655    0.5286655    0.5286655
0.5398693
## MAE        338.0400480 338.0400480 338.0400480 338.0400480
254.6791507
##          pred_metric pred_metric pred_metric pred_metric
pred_metric
## RMSE      482.8296811 482.8296811 482.8296811 568.6757233
568.6757233
## Rsquared    0.5398693    0.5398693    0.5398693    0.2851607
0.2851607
## MAE        254.6791507 254.6791507 254.6791507 436.9034926
436.9034926
##          pred_metric pred_metric
## RMSE      568.6757233 568.6757233
## Rsquared    0.2851607    0.2851607
## MAE        436.9034926 436.9034926

names_var <- c()
for (i in a) {
  for(j in b) {
    names_var <- append(names_var, paste(i, j))
  }
}

names_var

```

```
## [1] "Volume ~ x4StarReviews lm"
## [2] "Volume ~ x4StarReviews rf"
## [3] "Volume ~ x4StarReviews knn"
## [4] "Volume ~ x4StarReviews svmLinear"
## [5] "Volume ~. lm"
## [6] "Volume ~. rf"
## [7] "Volume ~. knn"
## [8] "Volume ~. svmLinear"
## [9] "Volume ~ PositiveServiceReview lm"
## [10] "Volume ~ PositiveServiceReview rf"
## [11] "Volume ~ PositiveServiceReview knn"
## [12] "Volume ~ PositiveServiceReview svmLinear"
```

```
colnames(compare_var_mod) <- names_var
```

```
compare_var_mod
```

```
##          Volume ~ x4StarReviews lm Volume ~ x4StarReviews rf
## RMSE          459.4407974          459.4407974
## Rsquared        0.5286655          0.5286655
## MAE            338.0400480          338.0400480
##          Volume ~ x4StarReviews knn Volume ~ x4StarReviews
svmLinear
## RMSE          459.4407974
459.4407974
## Rsquared        0.5286655
0.5286655
## MAE            338.0400480
338.0400480
##          Volume ~. lm Volume ~. rf Volume ~. knn Volume ~.
svmLinear
## RMSE          482.8296811  482.8296811  482.8296811
482.8296811
## Rsquared        0.5398693  0.5398693  0.5398693
0.5398693
## MAE            254.6791507  254.6791507  254.6791507
254.6791507
##          Volume ~ PositiveServiceReview lm
## RMSE          568.6757233
## Rsquared        0.2851607
## MAE            436.9034926
##          Volume ~ PositiveServiceReview rf
## RMSE          568.6757233
## Rsquared        0.2851607
## MAE            436.9034926
##          Volume ~ PositiveServiceReview knn
## RMSE          568.6757233
## Rsquared        0.2851607
## MAE            436.9034926
##          Volume ~ PositiveServiceReview svmLinear
```

```
## RMSE 568.6757233
## Rsquared 0.2851607
## MAE 436.9034926

compare_var_mod_melt <- melt(compare_var_mod, varnames =
c("metric", "model"))
compare_var_mod_melt <- as.data.frame(compare_var_mod_melt)
compare_var_mod_melt
```

##	metric	model
value		
## 1	RMSE	Volume ~ x4StarReviews lm
459.4407974		
## 2	Rsquared	Volume ~ x4StarReviews lm
0.5286655		
## 3	MAE	Volume ~ x4StarReviews lm
338.0400480		
## 4	RMSE	Volume ~ x4StarReviews rf
459.4407974		
## 5	Rsquared	Volume ~ x4StarReviews rf
0.5286655		
## 6	MAE	Volume ~ x4StarReviews rf
338.0400480		
## 7	RMSE	Volume ~ x4StarReviews knn
459.4407974		
## 8	Rsquared	Volume ~ x4StarReviews knn
0.5286655		
## 9	MAE	Volume ~ x4StarReviews knn
338.0400480		
## 10	RMSE	Volume ~ x4StarReviews svmLinear
459.4407974		
## 11	Rsquared	Volume ~ x4StarReviews svmLinear
0.5286655		
## 12	MAE	Volume ~ x4StarReviews svmLinear
338.0400480		
## 13	RMSE	Volume ~. lm
482.8296811		
## 14	Rsquared	Volume ~. lm
0.5398693		
## 15	MAE	Volume ~. lm
254.6791507		
## 16	RMSE	Volume ~. rf
482.8296811		
## 17	Rsquared	Volume ~. rf
0.5398693		
## 18	MAE	Volume ~. rf
254.6791507		
## 19	RMSE	Volume ~. knn
482.8296811		
## 20	Rsquared	Volume ~. knn
0.5398693		

```


## 21      MAE      Volume ~. knn
254.6791507
## 22      RMSE      Volume ~. svmLinear
482.8296811
## 23 Rsquared      Volume ~. svmLinear
0.5398693
## 24      MAE      Volume ~. svmLinear
254.6791507
## 25      RMSE      Volume ~ PositiveServiceReview lm
568.6757233
## 26 Rsquared      Volume ~ PositiveServiceReview lm
0.2851607
## 27      MAE      Volume ~ PositiveServiceReview lm
436.9034926
## 28      RMSE      Volume ~ PositiveServiceReview rf
568.6757233
## 29 Rsquared      Volume ~ PositiveServiceReview rf
0.2851607
## 30      MAE      Volume ~ PositiveServiceReview rf
436.9034926
## 31      RMSE      Volume ~ PositiveServiceReview knn
568.6757233
## 32 Rsquared      Volume ~ PositiveServiceReview knn
0.2851607
## 33      MAE      Volume ~ PositiveServiceReview knn
436.9034926
## 34      RMSE Volume ~ PositiveServiceReview svmLinear
568.6757233
## 35 Rsquared Volume ~ PositiveServiceReview svmLinear
0.2851607
## 36      MAE Volume ~ PositiveServiceReview svmLinear
436.9034926

```

```

ggplot(compare_var_mod_melt, aes(x=model,y=value)) + geom_col() +
facet_grid(metric~., scales="free")
+theme(axis.text=element_text(size=3),
        axis.title=element_text(size=14,face="bold"))

```



I only used RF, and KNN because from past results the SVM did not look like a good fit.

Error Analysis

```
ABSrf <- (List$testingSet$Volume - PredictionRandomForest)
```

```
RLTrf <- (ABSrf / List$testingSet$Volume)
```

```
ABSsvm <- (List$testingSet$Volume - svm.pred)
```

```
RLTsvm <- (ABSsvm / List$testingSet$Volume)
```

```
Absknn <- (List$testingSet$Volume - KnnPrediction)
```

```
RLTknn <- (Absknn / List$testingSet$Volume)
#abline(0, 0) # the horizon
```


```
Lol <- cbind(List$testingSet, ABSrf)
```

Random Forest Residuals

```
ggplot(Lol,
  aes(Lol$Volume, ABSrf)) +
```

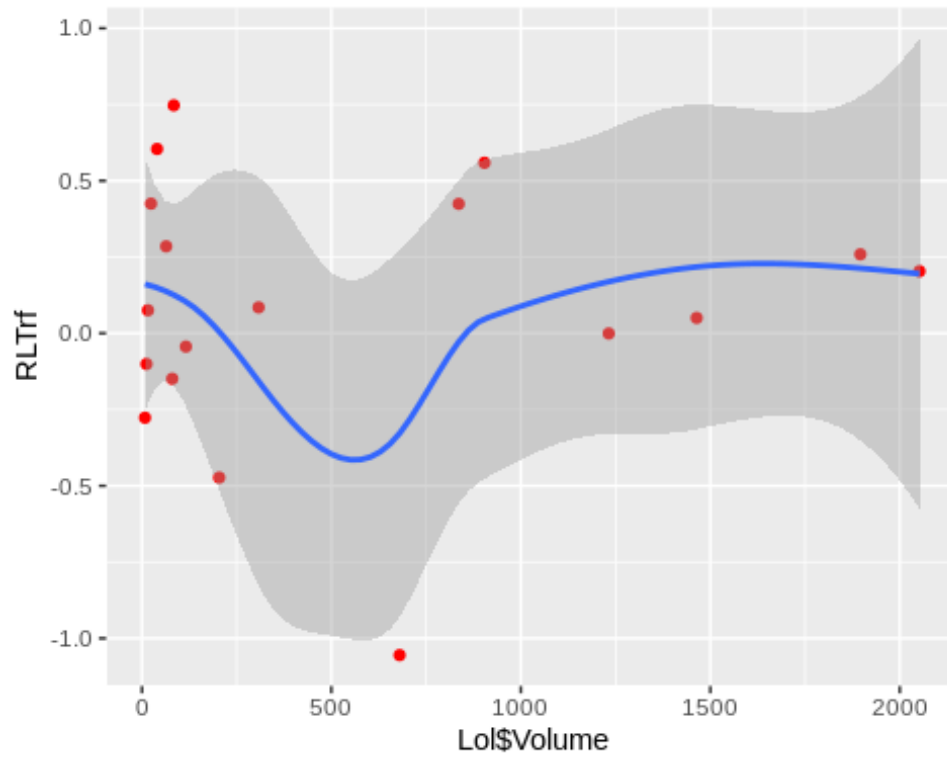
```
geom_point(color="red")+  
geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(Lol,  
  aes(Lol$Volume,RLTrf))+  
geom_point(color="red")+  
geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

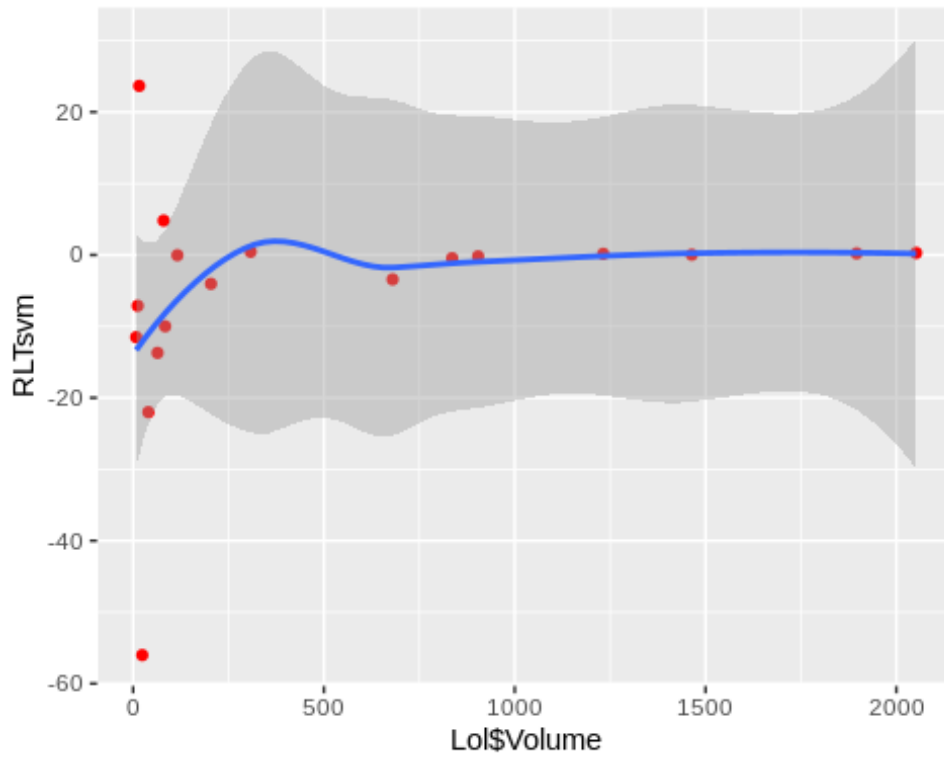


Residuals


```
ggplot(Lol,
  aes(Lol$Volume, ABSsvm)) +
  geom_point(color="red") +
  geom_smooth()
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



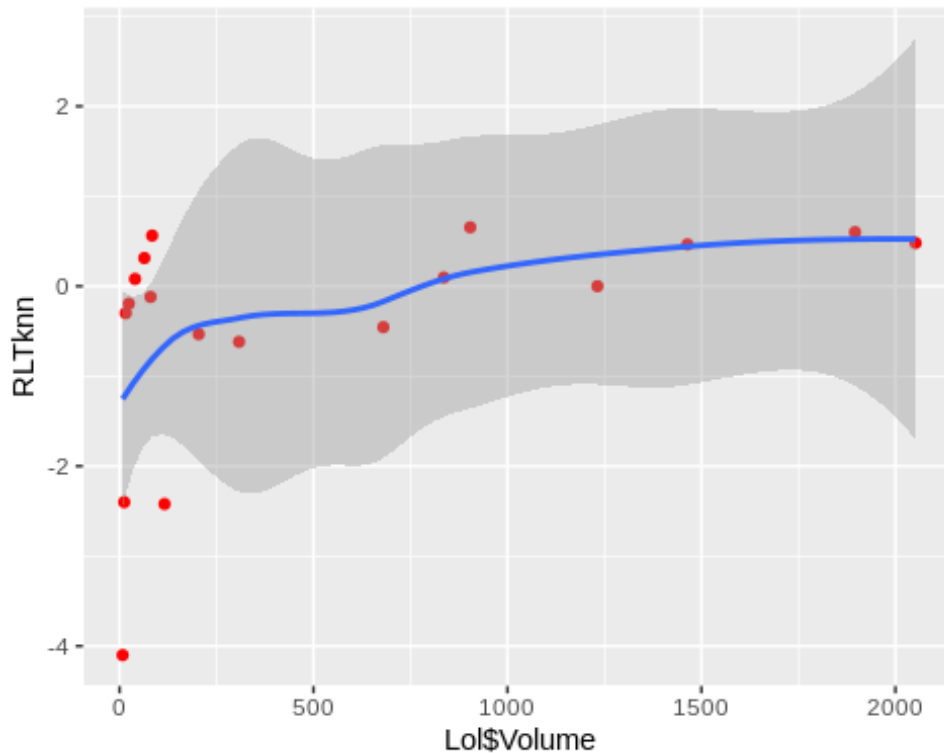
```
ggplot(Lol,  
       aes(Lol$Volume,RLTsvm))+  
  geom_point(color="red")+  
  geom_smooth()  
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

```
ggplot(Lol,  
       aes(Lol$Volume,Absknn))+  
  geom_point(color="red")+  
  geom_smooth()  
  
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(Lol,  
  aes(Lol$Volume,RLTknn))+  
  geom_point(color="red")+  
  geom_smooth()  
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Let's now apply the current models into the new product list and make a top 5 for most probably sold products in volume.

Prediction

The random forest was the one who gave us the best results, with both variables (ProductServiceReview,x4StarReviews).

```
NP <- read.csv(file =
"/home/zordo/Documents/Ubiquum/R-M2Task3/data/Npa.csv", header =
TRUE , sep = ',')
```

```
NP <- PPfunction(NP)
```

```
NewProductsVolume <- predict(ModelRandomForest, NP)
```

```
NP$Volume<-NewProductsVolume
```

We got the volume, now we need to calculate the profit, to see which products types we should invest on.

Profit = profit margin * Volume

```

Profit <- NP$ProfitMargin * NP$Volume

NP <- cbind(NP,Profit)

Top5 <- top_n(NP, 5, Profit)

Top5 <- cbind (Top5,sort(Top5$Profit))

Top5

##   ProductType.Accessories ProductType.Display
##   ProductType.ExtendedWarranty
## 1                0                0
## 0
## 2                0                0
## 0
## 3                0                0
## 0
## 4                0                0
## 0
## 5                0                0
## 0
##   ProductType.GameConsole ProductType.Laptop
##   ProductType.Netbook
## 1                0                0
## 1
## 2                0                0
## 0
## 3                0                0
## 0
## 4                1                0
## 0
## 5                1                0
## 0
##   ProductType.PC ProductType.Printer
##   ProductType.PrinterSupplies
## 1                0                0
## 0
## 2                0                0
## 0
## 3                0                0
## 0
## 4                0                0
## 0
## 5                0                0
## 0
##   ProductType.Smartphone ProductType.Software
##   ProductType.Tablet
## 1                0                0

```

```

0
## 2          0          0
1
## 3          0          0
1
## 4          0          0
0
## 5          0          0
0
##   ProductNum  Price x5StarReviews x4StarReviews x3StarReviews
## 1         180 329.00          312          112          28
## 2         186 629.00          296           66          30
## 3         187 199.00          943          437         224
## 4         199 249.99          462           97          25
## 5         307 425.00         1525          252          99
##   x2StarReviews x1StarReviews PositiveServiceReview
NegativeServiceReview
## 1           31           47           28
16
## 2           21           36           28
9
## 3          160          247           90
23
## 4           17           58           32
12
## 5           56           45           59
13
##   Recommendproduct BestSellersRank ShippingWeight ProductDepth
## 1             0.7          2699           4.6          10.17
## 2             0.8           34           3.0           7.31
## 3             0.8            1           0.9           5.40
## 4             0.8          115           8.4           6.20
## 5             0.9          215          20.0           8.50
##   ProductWidth ProductHeight ProfitMargin  Volume  Profit
## 1          7.28          0.95          0.09 1386.986 124.8287
## 2          9.50          0.37          0.10 1371.081 137.1081
## 3          7.60          0.40          0.20 1835.990 367.1979
## 4         13.20         13.20          0.09 1513.556 136.2200
## 5          6.00          1.75          0.18 1843.867 331.8960
##   sort(Top5$Profit)
## 1          124.8287
## 2          136.2200
## 3          137.1081
## 4          331.8960
## 5          367.1979

```

Our top 5 most profitable product types are

Number

Game Console

Game Console

Tablet

Tablet

NoteBook

Conclusion

All three models used are non-parametric, but before explaining what a non-parametric model is, I would like to explain what parametric models are. Parametric models are algorithms that simplify the function to a known form, and no matter how much data that's fed to the algorithm, the model won't change the quantity of parameters needed. All you need to know for predicting a future data value from the current state of the model are its parameters, EG : Linear regression with one variable, you have two parameters (coefficient and intercept). Knowing these parameters will enable you to predict new values. In a mathematical way :

$Y_i = B_0 + B_1X_1 + B_2X_2 + \dots + e_i$ Non-parametric models do not make any fixed assumptions about the form of the mapping, they are free to learn from the training data. The parameters are usually said to be infinite in dimension and so can express the characteristics in data much better than parametric models. E.g : KNN, makes predictions based on the k most similar training patterns for a new data instance. The method does not assume anything about the form of the mapping function other than patterns that are close are likely to have a similar output variable. In mathematical form **$Y_i = F(x_i) + e_i$** Where F can be any function, the data will decide what the function looks like. It will not provide the analytical expression but it will give you its graph given your data set. I find this picture very helpful in understanding how both types of models work, summing up in easy non-technical language, parametric models follow an equation while non-parametric models follow the data. One of the requirements to use a non-parametric model is to have a big set of data, we have 80 rows, 78 after cutting outliers and only 60 of them are for training. This is not a big set of data at all, I think it actually couldn't be any smaller than this. Also they are good methods when we don't

have any prior knowledge and not worry about choosing the right features. So using this models is not efficient at all, if we are using algorithms that “follow the data” and we have almost no data, it’s obvious why this approach is not efficient and why the errors are so big.